

Agent Factory Certification Quiz

Building Digital Full-Time Equivalents (FTEs) 90 Multiple Choice Questions with Answers & Rationales
January 2026 Edition

Q1. A company already uses predictive analytics to forecast churn. They now want AI that can take actions automatically (sending retention offers, updating CRM, scheduling follow-ups). Which AI wave does this transition represent?

- A. Predictive AI
- B. Generative AI
- C. **Agentic AI ✓**
- D. Analytical AI

Correct Answer: C

Rationale: Agentic AI moves beyond prediction/content to autonomous action loops.

Q2. Why does the curriculum emphasize “Code as the Universal Interface” rather than code as just software development?

- A. Code is faster than natural language
- B. **Code allows AI agents to interrogate real systems and data ✓**
- C. Code replaces the need for specs
- D. Code reduces token costs

Correct Answer: B

Rationale: Code lets agents query databases, call APIs, execute actions—interacting with reality.

Q3. Which characteristic most clearly distinguishes an AI Agent from a chatbot?

- A. Conversational ability
- B. Large language model usage
- C. **Autonomous goal pursuit with tool use ✓**
- D. Natural language understanding

Correct Answer: C

Rationale: Agents pursue goals, use tools, maintain state, and self-correct. Chatbots mainly respond to inputs.

Q4. In the Agent Factory worldview, what fundamentally changes when moving from tool-centric enterprises to agent-centric enterprises?

- A. SaaS tools are eliminated
- B. Humans stop working
- C. **Logic moves from people's heads into specs and skills ✓**
- D. APIs are no longer needed

Correct Answer: C

Rationale: Knowledge becomes reusable intellectual property encoded in specifications and skills.

Q5. Why is Markdown treated as a “programming language” in Agent Factory?

- A. Markdown executes faster
- B. Markdown is easier than Python
- C. **AI agents can reliably execute structured natural language ✓**
- D. Markdown replaces SDKs

Correct Answer: C

Rationale: Structured natural language in Markdown becomes executable instructions for AI agents.

Q6. Which of the following best represents the OODA loop used by General Agents?

- A. Prompt → Response → Save
- B. **Observe → Decide → Act → Correct ✓**
- C. Train → Deploy → Monitor
- D. Predict → Generate → Answer

Correct Answer: B

Rationale: The OODA loop (Observe–Orient–Decide–Act, extended with correction) enables reasoning, action, and self-improvement.

Q7. Why does the curriculum warn against calling Claude Code merely a “coding agent”?

- A. It cannot write code
- B. It only works in terminals
- C. **It solves business problems using code as a tool ✓**

D. It is slower than IDE-based agents

Correct Answer: C

Rationale: Code is the mechanism – the real purpose is solving broader business problems.

Q8. Which of the following best explains Agent Skills (SKILL.md)?

- A. Prompt templates reused across chats
- B. Fine-tuned model weights
- C. Deterministic, reusable expertise encoded as instructions + logic ✓
- D. MCP server configurations

Correct Answer: C

Rationale: Skills are standardized, reusable, deterministic packages of business expertise.

Q9. What is the primary purpose of MCP (Model Context Protocol)?

- A. Improve language generation quality
- B. Reduce hallucinations
- C. Connect agents to live external systems ✓
- D. Replace APIs

Correct Answer: C

Rationale: MCP serves as the standardized data and action bridge between agents and external tools/systems.

Q10. A customer service workflow must handle 5,000 identical tickets/day with strict compliance. Which agent type is most appropriate?

- A. General Agent
- B. Coding Agent
- C. Custom Agent via SDK ✓
- D. Chatbot

Correct Answer: C

Rationale: High volume + strict compliance needs = Custom Agent with strong guardrails and orchestration.

Q11. Why are General Agents considered expensive per task but valuable?

- A. They use larger models
- B. They require human supervision
- C. They handle non-routine, high-judgment work ✓
- D. They consume more memory

Correct Answer: C

Rationale: Their economic value lies in solving novel, ambiguous problems that rigid software cannot.

Q12. What is the key output of Spec-Driven Development?

- A. Prompts
- B. Conversations
- C. Executable specs that agents can implement ✓
- D. Training datasets

Correct Answer: C

Rationale: Specs serve as the "source code" that defines the logic and behavior of an agent.

Q13. Why does Agent Factory separate testing code (TDD) from testing reasoning (Agent Evals)?

- A. Reasoning cannot be tested
- B. Code accuracy guarantees reasoning accuracy
- C. Agents can be logically wrong even if code passes ✓
- D. Evals are cheaper than TDD

Correct Answer: C

Rationale: Reasoning quality must be evaluated semantically; code passing doesn't mean the logic was correct.

Q14. What role does the Golden Dataset play in Agent Evals?

- A. Model training
- B. Prompt tuning
- C. Pre-deployment examination of agent reasoning ✓
- D. Data labeling

Correct Answer: C

Rationale: It acts as a standardized "exam" to ensure the agent handles various scenarios correctly before going live.

Q15. Why does updating a SKILL.md require regression testing?

- A. To reduce token usage
- B. **To ensure previous competencies still perform correctly ✓**
- C. To retrain the model
- D. To update MCP schemas

Correct Answer: B

Rationale: Skills are reusable assets—changes to fix one path must not degrade past behavior in other paths.

Q16. Which element of a spec defines when a skill should activate?

- A. Identity
- B. Logic
- C. **Success Trigger ✓**
- D. Error Protocol

Correct Answer: C

Rationale: Triggers map environmental situations or inputs to specific skill activations.

Q17. Why does Agent Factory recommend external scripts (Python/Bash) for math and formatting?

- A. LLMs cannot calculate
- B. **Scripts reduce hallucination and increase determinism ✓**
- C. Scripts are cheaper than tokens
- D. SDKs require scripts

Correct Answer: B

Rationale: Deterministic logic belongs outside probabilistic models to ensure 100% accuracy in math/formatting.

Q18. Which deployment pattern minimizes risk when introducing agents into high-stakes workflows?

- A. Immediate full automation
- B. Batch processing
- C. **Shadow mode ✓**
- D. Fine-tuning first

Correct Answer: C

Rationale: Shadow mode allows validation of agent decisions against reality without actual consequences.

Q19. Why is observability critical in agent deployments?

- A. To debug SDK errors
- B. **To monitor accuracy, latency, and cost per task ✓**
- C. To reduce model size
- D. To comply with MCP

Correct Answer: B

Rationale: Agents are workers; their performance (ROI, speed, quality) must be measured like employees.

Q20. Why is selling a solution as a Digital FTE strategically superior to selling it as SaaS?

- A. It avoids compliance
- B. **It taps into HR/headcount budgets instead of IT budgets ✓**
- C. It lowers infrastructure costs
- D. It increases token usage

Correct Answer: B

Rationale: Headcount budgets are historically larger and value-based compared to fixed IT/software budgets.

Q21. What economic comparison most convinces executives to approve a Digital FTE?

- A. Token cost vs salary
- B. API latency vs human speed
- C. **Cost per task reduction (~85–90%) ✓**
- D. Model benchmark scores

Correct Answer: C

Rationale: Cost-per-task framing aligns directly with corporate ROI and efficiency thinking.

Q22. Which monetization model is most suitable for regulated industries that require data to remain in-house?

- A. Digital FTE Subscription
- B. Success Fee
- C. Skill Marketplace
- D. License Model ✓

Correct Answer: D

Rationale: Licensing sells the "recipe" (specs/skills) while allowing the data to remain inside the client's secure perimeter.

Q23. Why does the License model treat SKILL.md as the primary asset?

- A. It is easier to copy
- B. It contains codified business logic ✓
- C. It replaces models
- D. It reduces infrastructure costs

Correct Answer: B

Rationale: Skills encode proprietary business expertise, which is the actual value being licensed.

Q24. What was the key success factor in the Digital SDR case study?

- A. Better model selection
- B. Prompt engineering
- C. Codifying sales expertise into a reusable Agent Skill ✓
- D. Hiring more SDRs

Correct Answer: C

Rationale: Turning human expertise into a reusable skill asset allows for massive, consistent scaling.

Q25. Why is multi-tenancy essential for scaling Digital FTEs?

- A. It improves reasoning
- B. It enables isolated data handling across enterprises ✓
- C. It reduces prompt size
- D. It simplifies specs

Correct Answer: B

Rationale: Enterprise trust requires strict isolation between different customers' data and contexts.

Q26. What problem does cloud-native auto-scaling directly solve in Agent Factory?

- A. Prompt drift
- B. Model hallucinations
- C. Supporting sudden growth in customers without downtime ✓
- D. Skill versioning

Correct Answer: C

Rationale: Auto-scaling allows the factory to spin up as many "digital employees" as needed to meet demand.

Q27. Why is vendor lock-in considered a major strategic risk?

- A. Models evolve slowly
- B. Token prices increase
- C. Agent skills become non-portable ✓
- D. Specs stop working

Correct Answer: C

Rationale: Lock-in prevents the reuse of your codified IP across different platforms or newer models.

Q28. What is the Scaling Paradox described in the curriculum?

- A. AI cannot scale
- B. Infrastructure is expensive
- C. Human effort must be decoupled from business growth ✓
- D. Models plateau in intelligence

Correct Answer: C

Rationale: To achieve "Unicorn" scale, revenue must grow non-linearly compared to human headcount.

Q29. Why is OpenAI Apps described as the "App Store moment" for agents?

- A. Lower token prices

- A. Built-in authentication
- C. Instant global distribution without enterprise sales cycles ✓
- D. Better SDKs

Correct Answer: C

Rationale: Distribution, not just technology, is what unlocks global scale and rapid adoption.

Q30. According to the Agent Factory "Golden Rule," what ultimately determines agent quality?

- A. Model size
- B. Prompt creativity
- C. The clarity and completeness of the spec ✓
- D. SDK choice

Correct Answer: C

Rationale: "Your spec is your source code." The blueprint determines the outcome.

Q31. An enterprise wants an agent to analyze financial data only when a CSV is uploaded and the phrase "quarterly audit" is used. Which spec component governs this behavior?

- A. Identity
- B. Logic
- C. Success Trigger ✓
- D. Output Standard

Correct Answer: C

Rationale: Success triggers define the specific context and conditions under which a skill activates.

Q32. Why does Agent Factory treat SKILL.md as a form of Intellectual Property (IP)?

- A. It contains proprietary model weights
- B. It represents reusable, codified business expertise ✓
- C. It reduces cloud costs
- D. It replaces MCP servers

Correct Answer: B

Rationale: SKILL.md encodes domain expertise into a deterministic digital asset a company can own.

Q33. Which failure mode is most likely if edge cases are not documented in the spec?

- A. Increased token cost
- B. Agent hallucinations in rare scenarios ✓
- C. Slower deployment
- D. SDK incompatibility

Correct Answer: B

Rationale: Without clear instructions for edge cases, agents "guess" based on probability (hallucination).

Q34. What is the primary architectural reason MCP is preferred over custom point-to-point integrations?

- A. Lower latency
- B. Easier prompt writing
- C. Standardized, reusable connectivity across agents ✓
- D. Reduced hallucinations

Correct Answer: C

Rationale: MCP provides a universal standard for connecting any agent to any external system.

Q35. Why are General Agents ideal for the "Builder" role in the Agent Factory?

- A. They are cheaper to run
- B. They can fine-tune models
- C. They can reason across domains and generate new skills or agents ✓
- D. They replace SDKs

Correct Answer: C

Rationale: Builders must reason, plan, and create other assets, requiring general intelligence.

Q36. Which part of the Agent Factory workflow converts human intent into production assets?

- A. MCP Servers
- B. The Builder (Claude Code) ✓
- C. The Custom Agent SDK

D. Cloud Infrastructure

Correct Answer: B

Rationale: The Builder interprets high-level specs and generates executable skills and agents.

Q37. Why is manual prompting considered inferior to Agent Skills for enterprise use?

- A. Prompts are too short
- B. Prompts cannot call tools
- C. **Prompts are disposable and non-deterministic ✓**
- D. Prompts increase latency

Correct Answer: C

Rationale: Prompts lack the reliability, determinism, and version control required for enterprise operations.

Q38. Which metric best evaluates whether an agent is enterprise-ready?

- A. Model benchmark score
- B. Token usage
- C. **Accuracy over a golden dataset with regression stability ✓**
- D. Prompt length

Correct Answer: C

Rationale: Enterprises require consistency and the ability to prove that updates don't break existing logic.

Q39. Why does Agent Factory recommend human-in-the-loop (HITL) for certain workflows?

- A. Agents are unreliable
- B. **Compliance and high error-cost scenarios require oversight ✓**
- C. It reduces cloud cost
- D. It improves reasoning speed

Correct Answer: B

Rationale: HITL is a risk-management strategy for high-stakes decisions, not a sign of agent weakness.

Q40. An agent must never approve refunds over \$500. Where should this rule live?

- A. MCP server
- B. Success Trigger
- C. **Logic section ("Never list") ✓**
- D. Output template

Correct Answer: C

Rationale: Hard constraints and business boundaries must be clearly codified in the agent's logic.

Q41. Why does the curriculum argue that most AI agent failures are organizational, not technical?

- A. Models are already perfect
- B. Infrastructure is mature
- C. **Change management, unclear KPIs, and adoption issues dominate ✓**
- D. SDKs are unstable

Correct Answer: C

Rationale: Technology is rarely the bottleneck; human process, buy-in, and measurement are the real challenges.

Q42. Which KPI is least useful when evaluating a Digital FTE?

- A. Cost per task
- B. Accuracy over time
- C. **Number of tokens consumed ✓**
- D. Time saved

Correct Answer: C

Rationale: Token consumption is a technical byproduct; business value is found in cost, time, and accuracy.

Q43. Why does the Digital FTE model unlock higher pricing power?

- A. AI models are expensive
- B. **Customers compare it to employee salaries, not software licenses ✓**
- C. It requires custom infrastructure
- D. It uses premium SDKs

Correct Answer: B

Rationale: Anchoring against a human salary (e.g., \$60k) allows for much higher pricing than a \$50/mo SaaS seat.

Q44. Which monetization model best aligns risk-sharing between vendor and client?

- A. License
- B. Digital FTE Subscription
- C. Skill Marketplace
- D. **Success Fee ✓**

Correct Answer: D

Rationale: Success fees ensure the vendor is only paid when the agent delivers a specific, valuable outcome.

Q45. Why is shadow mode strategically recommended before full automation?

- A. To train the model
- B. To improve prompts
- C. **To validate performance without business impact ✓**
- D. To reduce compute usage

Correct Answer: C

Rationale: Shadow mode derisks deployment by allowing agents to "practice" on live data without acting on it.

Q46. What allowed the Digital SDR agent to outperform human teams so dramatically?

- A. Better data
- B. Larger models
- C. **Instant scaling + deterministic follow-up logic ✓**
- D. Cheaper tokens

Correct Answer: C

Rationale: Agents provide the perfect combination of infinite scale and zero "forgetfulness" in follow-ups.

Q47. Why is agent duplication considered exponential scaling?

- A. Models get smarter over time
- B. **One skill can be instantiated infinitely at near-zero marginal cost ✓**
- C. Cloud providers optimize GPUs
- D. Token prices decrease

Correct Answer: B

Rationale: Digital assets allow for non-linear growth because the marginal cost of a new "worker" is negligible.

Q48. Which role is most critical for ensuring an agent actually delivers business value?

- A. MLOps Engineer
- B. Security Lead
- C. **AI Product Owner ✓**
- D. SDK Maintainer

Correct Answer: C

Rationale: The AI Product Owner bridges the gap between technical capability and business KPIs.

Q49. Why is multi-tenancy non-negotiable for enterprise agent platforms?

- A. Faster inference
- B. Cheaper hosting
- C. **Secure isolation of each customer's data and context ✓**
- D. Easier debugging

Correct Answer: C

Rationale: Security and privacy (data isolation) are the first requirements for any enterprise-grade deployment.

Q50. What does the Scaling Paradox force AI startups to optimize for?

- A. Better models
- B. More features
- C. **Distribution and automation over human effort ✓**
- D. Cheaper GPUs

Correct Answer: C

Rationale: To reach massive scale, the business must decouple its growth from its human employee count.

Q51. Why is OpenAI Apps strategically important for Agent Factory businesses?

- A. Lower token prices
- B. Built-in compliance
- C. Instant access to millions of users and enterprises ✓
- D. Better reasoning models

Correct Answer: C

Rationale: Distribution is the "App Store" moment—getting the product in front of buyers without long sales cycles.

Q52. Which scenario should NOT be automated according to the curriculum?

- A. Invoice reconciliation
- B. Lead qualification
- C. High-stakes legal judgment without review ✓
- D. Customer support triage

Correct Answer: C

Rationale: Irreversible decisions with massive legal or ethical weight require human accountability.

Q53. What is the primary goal of Agent Eval from a buyer's perspective?

- A. Reduce hallucinations
- B. Ensure the agent passes an "exam" before being trusted ✓
- C. Improve token efficiency
- D. Fine-tune prompts

Correct Answer: B

Rationale: Buyers need a "certification of quality" to trust a digital employee with their business process.

Q54. Why does the curriculum emphasize "Outcome over Usage" pricing?

- A. Easier billing
- B. Lower infrastructure costs
- C. Aligns AI value with business results ✓
- D. Reduces prompt complexity

Correct Answer: C

Rationale: Businesses don't want to buy "AI usage"; they want to buy completed work.

Q55. Which component ensures enterprise compliance auditing?

- A. Specs
- B. MCP
- C. Immutable audit logs ✓
- D. Agent Skills

Correct Answer: C

Rationale: Compliance requires a permanent, unchangeable record of every action the agent took.

Q56. Why does Agent Factory recommend 1 Agent Engineer per 3–5 production agents?

- A. To reduce salary cost
- B. To handle monitoring, updates, and edge cases effectively ✓
- C. Because SDKs require it
- D. Because agents degrade quickly

Correct Answer: B

Rationale: Production agents require human stewardship to maintain accuracy and handle environmental changes.

Q57. What differentiates a Custom Agent from a General Agent at scale?

- A. Intelligence level
- B. Guardrails and orchestration control ✓
- C. Model size
- D. Token limits

Correct Answer: B

Rationale: Custom agents provide the orchestration needed for predictable, safe results in specific domains.

Q58. Why is spec clarity more important than model choice?

- A. Models are identical

- B. Specs determine behavior, constraints, and success ✓
- C. Specs reduce compute cost
- D. Specs replace SDKs

Correct Answer: B

Rationale: A powerful model with a bad spec will produce a high-quality failure.

Q59. Which statement best captures the Agent Factory thesis?

- A. AI will replace all workers
- B. Software is obsolete
- C. **The future is manufacturing digital employees, not apps ✓**
- D. Models will self-improve endlessly

Correct Answer: C

Rationale: The transition is from "tools for humans" to "autonomous digital labor."

Q60. What ultimately allows an Agent Factory business to reach unicorn scale?

- A. Best-in-class models
- B. Custom infrastructure
- C. **Reusable specs, skills, and automated distribution ✓**
- D. Large sales teams

Correct Answer: C

Rationale: Scalability is driven by reusable codified IP and frictionless global delivery.

Q61. Why does the curriculum frame the future enterprise not around software tools but around digital teammates?

- A. Tools are too expensive
- B. Agents reduce UI complexity
- C. **Work is outcome-driven, not interaction-driven ✓**
- D. SaaS innovation has stopped

Correct Answer: C

Rationale: Enterprises care about the work getting done, not the dashboard used to do it.

Q62. What fundamentally distinguishes a Digital FTE from traditional automation (RPA)?

- A. Uses LLMs
- B. Is cheaper
- C. **Encodes judgment, not just rules ✓**
- D. Runs in the cloud

Correct Answer: C

Rationale: Digital FTEs handle non-routine tasks that require context and decision-making.

Q63. Why is agent autonomy economically valuable?

- A. It increases token usage
- B. **It reduces the need for human supervision ✓**
- C. It improves model benchmarks
- D. It replaces MCP

Correct Answer: B

Rationale: Less human supervision equals a lower cost per task and higher scaling potential.

Q64. Which pricing argument resonates most with a CFO?

- A. "This uses GPT-5"
- B. "It has 99.9% uptime"
- C. **"This replaces \$80k of annual labor for \$12k" ✓**
- D. "It's built with MCP"

Correct Answer: C

Rationale: CFOs view software through the lens of labor arbitrage and bottom-line impact.

Q65. Why does the curriculum emphasize cost per task over cost per token?

- A. Tokens fluctuate
- B. Tokens are confusing
- C. **Cost per task aligns directly with business value ✓**
- D. Tokens are deprecated

Correct Answer: C

Rationale: Business owners need to know the cost of the outcome, not the cost of the computation.

Q66. What makes specs a defensible competitive moat?

- A. They are proprietary model weights
- B. **They encode unique organizational knowledge ✓**
- C. They are hard to write
- D. They reduce cloud bills

Correct Answer: B

Rationale: Codified business logic and domain expertise are harder for competitors to replicate than buying model access.

Q67. Why does the Agent Factory strategy discourage "automation everywhere"?

- A. AI is unreliable
- B. Automation is expensive
- C. **Poorly chosen automation destroys trust and ROI ✓**
- D. Humans resist automation

Correct Answer: C

Rationale: Strategic focus on high-impact areas is better than "sprinkling AI" on low-value tasks.

Q68. What is the key risk of automating a process with undefined success criteria?

- A. Higher latency
- B. Model hallucinations
- C. **Inability to validate performance ✓**
- D. Increased cloud spend

Correct Answer: C

Rationale: If you don't know what "good" looks like, you can't tell if the agent is failing or succeeding.

Q69. Why is human review within SLA a critical decision checkpoint?

- A. To train agents
- B. To satisfy auditors
- C. **To ensure errors can be corrected before impact ✓**
- D. To reduce compute usage

Correct Answer: C

Rationale: Human review for edge cases or high-risk actions prevents catastrophic failures.

Q70. Which statement best reflects the Future of Work described in the curriculum?

- A. Humans will be replaced
- B. AI will work independently
- C. **Humans, agents, and robots will collaborate ✓**
- D. Knowledge work will disappear

Correct Answer: C

Rationale: The future is an ecosystem of specialized agents managed by high-leverage humans.

Q71. Why is change management essential for agent adoption?

- A. Agents require training
- B. **Teams fear job loss and resist adoption ✓**
- C. SDKs are complex
- D. Models drift

Correct Answer: B

Rationale: Technical success is impossible without addressing the human/cultural impact of AI labor.

Q72. Which executive concern does the Digital FTE framing directly address?

- A. Model accuracy
- B. Developer productivity
- C. **Headcount efficiency ✓**
- D. Token optimization

Correct Answer: C

Rationale: Digital FTEs are a solution for doing more work without growing human headcount linearly.

Q73. Why does the curriculum treat distribution as more important than model quality at scale?

- A. Models are equal
- B. Distribution determines who captures value ✓**
- C. Models stop improving
- D. Distribution reduces cost

Correct Answer: B

Rationale: Value is captured by the player who controls the access to the end-user/enterprise.

Q74. What problem does the OpenAI Apps ecosystem solve for startups?

- A. Model training
- B. Compliance
- C. Enterprise sales friction ✓**
- D. Prompt engineering

Correct Answer: C

Rationale: It allows builders to bypass the "18-month sales cycle" by putting agents directly in front of users.

Q75. Why is cloud-native infrastructure critical to Agent Factory economics?

- A. Better AI reasoning
- B. Manual scaling
- C. Pay-for-use compute with elastic growth ✓**
- D. Simpler SDKs

Correct Answer: C

Rationale: Elastic infrastructure ensures margins remain high by matching costs directly to customer demand.

Q76. Which factor most strongly influences enterprise trust in agents?

- A. Model provider brand
- B. UI design
- C. Security, auditability, and compliance ✓**
- D. Prompt quality

Correct Answer: C

Rationale: Enterprises prioritize safety and legal defensibility over fancy features.

Q77. Why is immutable audit logging non-negotiable?

- A. Debugging
- B. Performance tuning
- C. Legal defensibility and traceability ✓**
- D. Cost optimization

Correct Answer: C

Rationale: A permanent record is required to prove what happened in the event of a dispute or error.

Q78. What distinguishes an AI product company from an Agent Factory company?

- A. Model ownership
- B. UI focus
- C. Manufacturing reusable intelligence at scale ✓**
- D. Cloud provider choice

Correct Answer: C

Rationale: Agent Factories focus on the systematic manufacturing of autonomous labor assets.

Q79. Why does the curriculum emphasize spec literacy as a core future skill?

- A. Fewer developers needed
- B. Specs replace coding entirely
- C. High-leverage humans will define outcomes, not write syntax ✓**
- D. Specs reduce bugs

Correct Answer: C

Rationale: Humans are moving up the stack from "typing code" to "orchestrating outcomes."

Q80. What limits a company's ability to scale Digital FTEs globally?

- A. Model intelligence
- B. Cloud capacity
- C. Manual customer support
- D. **Spec and skill reusability ✓**

Correct Answer: D

Rationale: If every agent requires custom work, you are a services company, not a scalable factory.

Q81. Why does Agent Factory discourage vendor-specific logic inside skills?

- A. Harder to write
- B. Higher cost
- C. **Reduced portability and long-term leverage ✓**
- D. Worse performance

Correct Answer: C

Rationale: Skills should be generic enough to run on any model or platform to protect your IP.

Q82. What is the ultimate unit of value sold in the Agent Factory model?

- A. Tokens
- B. Models
- C. Software licenses
- D. **Codified labor ✓**

Correct Answer: D

Rationale: You aren't selling software; you are selling the work that the software performs.

Q83. Why is agent cloning strategically powerful?

- A. Faster inference
- B. Better reasoning
- C. **Zero-marginal-cost scaling of expertise ✓**
- D. Improved UX

Correct Answer: C

Rationale: Cloning allows a company to deploy a "top performer" agent to thousands of customers instantly.

Q84. Which signal best indicates agent maturity?

- A. Number of prompts
- B. Model version
- C. **Stable accuracy across updates ✓**
- D. SDK choice

Correct Answer: C

Rationale: Stability is the mark of a production-ready, reliable system.

Q85. Why does the curriculum warn against optimizing too early for cost?

- A. Costs don't matter
- B. Models are cheap
- C. **Premature optimization can block value discovery ✓**
- D. Cloud providers change pricing

Correct Answer: C

Rationale: Focus on proving the value and accuracy first; compute costs can be optimized once the ROI is clear.

Q86. Which mindset shift defines the developer-as-orchestrator role?

- A. Writing faster code
- B. Managing prompts
- C. **Designing systems of agents, skills, and specs ✓**
- D. Debugging models

Correct Answer: C

Rationale: The role moves from building features to building systems of autonomous workers.

Q87. Why is enterprise AI adoption accelerating faster than previous tech waves?

- A. Better marketing
- B. Cheaper hardware
- C. **Software is disrupting itself recursively ✓**

D. Regulations are lighter

Correct Answer: C

Rationale: AI tools are helping developers build AI agents, creating a positive feedback loop for growth.

Q88. What determines whether a Digital FTE is approved without extended debate?

- A. Model benchmarks
- B. Brand reputation
- C. **Clear 3–6 month payback period ✓**
- D. Technical architecture

Correct Answer: C

Rationale: Financial incentives (ROI) are the most powerful drivers of enterprise adoption.

Q89. Which statement best summarizes long-term Agent Factory advantage?

- A. Best prompts
- B. Cheapest tokens
- C. **Compounding intelligence assets ✓**
- D. Largest team

Correct Answer: C

Rationale: Advantage comes from accumulating a library of reusable, high-performance business skills.

Q90. What is the core promise of the Agent Factory vision?

- A. Fully autonomous AI
- B. No human oversight
- C. **Turning knowledge into scalable digital labor ✓**
- D. Replacing SaaS entirely

Correct Answer: C

Rationale: Knowledge is the input; digital labor (FTEs) is the industrial-grade output.